

PySpark on Dataproc

Apache Spark is an open source all-in-one data processing engine for big data workflows. It allows high performance batch processing, real-time streaming, SQL-querying, interactive analytics, data science and machine learning. It is written in Scala and supports variety of modern languages involving Python, Java, R and SQL. PySpark is an API to deal with Spark works in Python.

PySpark can be run locally or on cloud with variety of architectures that would be time consuming or risky while operating. Dataproc service of the GCP makes easier to handle the infrastructure for ETL tasks and data science as well as for Apache Spark in a fully managed and highly scalable manner.

Step0: Prerequisites

To follow this tutorial, one needs to have an active GCP account with activated Dataproc, BigQuery and Compute Engine APIs. Moreover, a dataset on BigQuery was prepared with dbt in previous tutorial for the project assignment of Data Engineering Zoomcamp by DataTalks.Club. It is also necessary to install and initialize Google Cloud CLI.

Step1: Creating a Dataproc cluster

It is a common practice to handle Apache Spark workloads with a flexible infrastructure that can scale immediately when needed. Hence, Dataproc builds them on clusters. So, the first step is to build one by clicking on *Create Cluster* button on DataProc page and select the option for running on a compute engine (the other option is for power users who want to use Kubernetes).

The screenshot displays the Google Cloud Platform Dataproc console. On the left, a sidebar shows navigation options like Clusters, Jobs, Workflows, and Autoscailing policies. The main area shows the 'Cluster details' for a cluster named 'cluster-31d1'. It includes fields for Name, Cluster UID, Type (Dataproc Cluster), and Status (Running). Below this, there are tabs for MONITORING, JOBS, VM INSTANCES, CONFIGURATION, and WEB INTERFACES. The MONITORING tab is active, showing YARN memory usage graphs for 'YARN memory' and 'YARN pending memory'. A 'Request to create cluster cluster-31d1 submitted' message is visible at the bottom. On the right, the 'Submit a job' form is open, with fields for Job ID, Job type (Hadoop), Main class or jar, Jar files, Archive files, Arguments, and Max restarts per hour. There are also sections for Properties and Labels.

- It is better to build your cluster on location where your BigQuery dataset is hosted.
- Even a single (1 master, 0 workers) node cluster is satisfactory for this simple tutorial.
- It would be helpful to involve Jupyter Notebook (and also Docker maybe) as optional components
- Default values for the rest of the settings are acceptable.

Step2: Transferring ML code

Having an up and running Dataproc cluster, you will have 2 more buckets called:

- dataproc-staging-...
- dataproc-temp-...

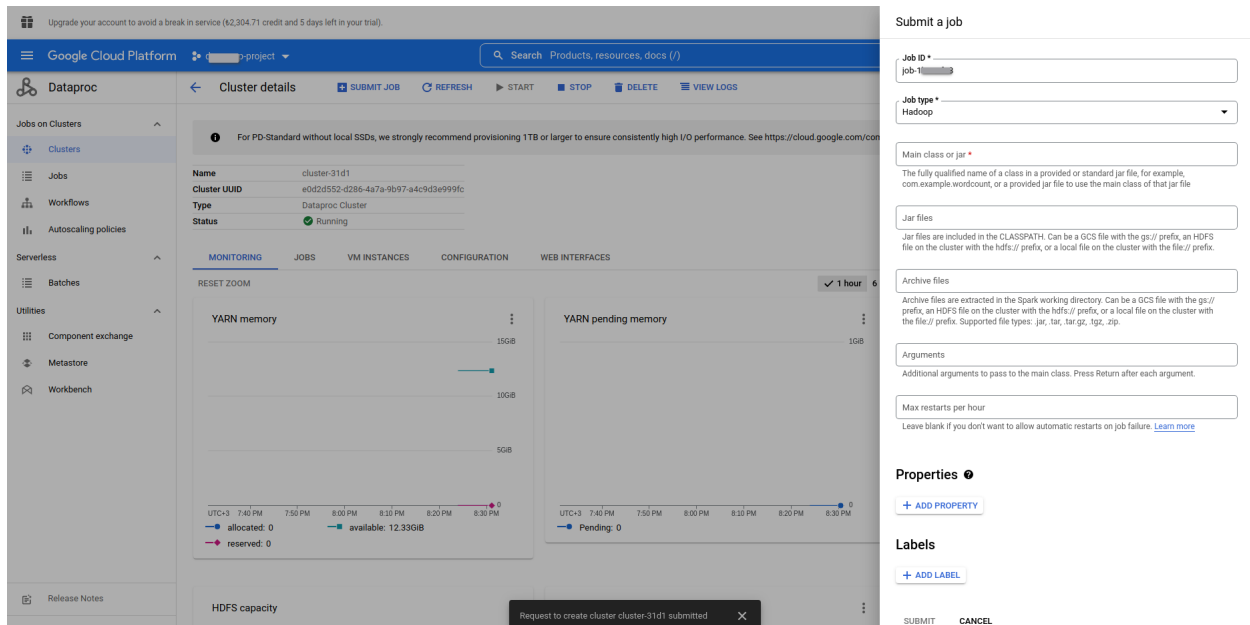
The next step is to transfer the machine learning code to the staging bucket using:

```
gsutil cp your_ml_code.py gs://your-bucket-name
```

Then, navigate through your code file on GCP bucket to copy the URL.

Step3: Submitting the ML task

Jump in to the Dataproc page on GCP and click on *Submit Job* button that will open a window:



Here, set:

- Job type as PySpark
- Ling for your ML code copied above (take care of the notation: `gs://prefix...`)
- Attach additional python file links if you have
- Insert `gs://spark-lib/bigquery/spark-bigquery-latest_2.12.jar` to involve required jar files.
- Submit the job

Step: Results

While the job is executed, the logs can be seen on Output tab where the results are also displayed.

Output

LINE WRAP: OFF

```
22/04/24 10:47:56 INFO org.sparkproject.jetty.server.AbstractConnector: Started ServerConnector@1e7092d9(HTTP/1.1, (http/1.1)){0.0.0.0:3437}
22/04/24 10:47:57 INFO org.apache.hadoop.yarn.client.RMProxy: Connecting to ResourceManager at cluster-cf02-m/10.172.0.5:8032
22/04/24 10:47:57 INFO org.apache.hadoop.yarn.client.AHSProxy: Connecting to Application History server at cluster-cf02-m/10.172.0.5:10200
22/04/24 10:47:58 INFO org.apache.hadoop.conf.Configuration: resource-types.xml not found
22/04/24 10:47:58 INFO org.apache.hadoop.yarn.util.resource.ResourceUtils: Unable to find 'resource-types.xml'.
22/04/24 10:47:59 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl: Submitted application application-1650745891311_0002
22/04/24 10:48:00 INFO org.apache.hadoop.yarn.client.RMProxy: Connecting to ResourceManager at cluster-cf02-m/10.172.0.5:8030
22/04/24 10:48:03 INFO com.google.cloud.hadoop.repackaged.gcs.com.google.cloud.hadoop.gcsio.GoogleCloudStorageImpl: Ignoring exception of type GoogleJsonResponseException; ve
spark initiated
data initiated
query initiated
22/04/24 10:48:12 INFO com.google.cloud.spark.bigquery.direct.DirectBigQueryRelation: Querying table dezcamp-project-57000.dbt_ckeskin.unified_all, parameters sent from Spark
22/04/24 10:48:12 INFO com.google.cloud.spark.bigquery.direct.DirectBigQueryRelation: Going to read from dezcamp-project-57000.dbt_ckeskin.unified_all columns=[measured_on, s
22/04/24 10:48:14 INFO com.google.cloud.spark.bigquery.direct.DirectBigQueryRelation: Created read session for table 'dezcamp-project-57000.dbt_ckeskin.unified_all': projects,
modeling initiated
22/04/24 10:48:35 WARN com.github.fommil.netlib.BLAS: Failed to load implementation from: com.github.fommil.netlib.NativeSystemBLAS
22/04/24 10:48:35 WARN com.github.fommil.netlib.BLAS: Failed to load implementation from: com.github.fommil.netlib.NativeRefBLAS
22/04/24 10:48:35 WARN com.github.fommil.netlib.LAPACK: Failed to load implementation from: com.github.fommil.netlib.NativeSystemLAPACK
22/04/24 10:48:35 WARN com.github.fommil.netlib.LAPACK: Failed to load implementation from: com.github.fommil.netlib.NativeRefLAPACK
Coefficients:[-6.930822923653599,0.12750885113789484,-0.0083132845228174,0.01661239265660598]
Intercept:9920.798478239238
R^2:0.6683789368000524
+-----+
| residuals |
+-----+
|-10.299100124922916|
|-10.099094696397515|
|-10.209488718464854|
|-10.0855923197214|
|-10.315523274386578|
|-10.258110201744785|
|-14.724515256695668|
|-10.168627714632748|
|-16.32527029392986|
|-10.460828055571255|
|-10.557646890808538|
|-10.31188104564717|
|-38.52232891495805|
|-10.577280619445446|
|-10.889632239406637|
|-10.174125362365885|
|-10.23137435659919|
|-10.273940420183862|
|-10.307178858989573|
|-10.372869638389602|
+-----+
only showing top 20 rows

completed
22/04/24 10:48:37 INFO org.sparkproject.jetty.server.AbstractConnector: Stopped Spark@1ef692d9(HTTP/1.1, (http/1.1)){0.0.0.0:0}
```

Output is complete

Moreover, while the cluster is running a task, it is possible to observ cluster workload and performance on *Job Details*

